

# Dimensioning

CGS - Hugo Miranda

2021



**Ciências  
ULisboa**

Faculdade  
de Ciências  
da Universidade  
de Lisboa

# Motivation

- Understand how to answer question such as:
  - Why is (will be) the system slow?
    - What is (are) the bottleneck(s)
  - Will the system handle the extra load?
    - How many replicas should I buy?
  - How does load flow from one component to the other?
  - What load to expect?

# Steps for Understanding Performance

- 1 Make a model
  - Divide into components
- 2 Populate with performance metrics
  - Estimates
  - Observed values
  - Hardware/software vendors
- 3 Analyze the model
  - Simulation
  - Mathematical analysis
- 4 Tune the model, restart

# Workload

**Workload** The amount of work per unit of time encountered by the **component**

- Unit: work/time (packets/s)
- Not necessarily the same for different components
  - 1 HTTP request
    - Multiple IP packets at the firewall
    - Even more frames (e.g. ARP)

# Capacity

**Capacity** The maximum possible throughput of the component

- Unit: work/time

# Response Time

**Response Time** The amount of time it takes for the component to respond to a request

- Unit: time
- Response time is not constant
  - Depends on the request
  - Cached web page vs DB query

# Requests Outcome

**Loss Rate** The fraction of requests that do not receive a response or receive an erroneous response

- Unit: %

**Throughput** (goodput) The amount of work per unit of time that receives a normal response through the component.

- Unit: work/time

# Utilization

**Utilization** The fraction of time the component is busy

- Unit: %
- Rule of thumb: aim to 30%



# Lets Forget Computers

<https://www.youtube.com/watch?v=MVm1KcrHM6s>  
Excerpt of "The Soup Nazi" (Seinfeld, S07E06, 1995)

# Our Restaurant Model

(Minutes as the unit of time)

## Operations table

Order	Operation	Time(m)	Component
1	Take order	10	Waiter
2	Cook meal	20	Chef
3	Eat	30	Table
4	Pay	5	Waiter
		65	

# Our Restaurant Model

## Resources

Component	Instances
Waiter	3
Chef	5
Table	20

# Graphical Representation



# Metrics Applied

Assume we have 10 clients/h

- Q: What is the **throughput** of the system?
  - A: If everything goes well,  $10\text{clients}/h = \frac{1}{6} \text{ clients}/m$
- Q: Is the **workload** equal for all components?
  - A: Throughput is the same for all components ( $\frac{1}{6}\text{clients}/m$ ) but components have a different number of instances
- Q: What is the **response time**?
  - A: See the table above:
    - For a meal: 65m
    - For a Chef: 20m

# Metrics Applied

Is the system properly sized? (or can it handle the workload?  
or Is the utilization  $< 1$ ?)

## Utilization Law

The average utilization of any component in the computer system is the throughput of the system multiplied by the amount of service time each request requires at that component

# Utilization Law In Practice

Component	Equation	Utilization
Waiter	$\frac{1}{6} \times \frac{10+5}{3}$	.83
Chef	$\frac{1}{6} \times \frac{20}{5}$	.66
Table	$\frac{1}{6} \times \frac{30}{20}$	.25

## ■ Notes:

- $\frac{1}{6}$  = throughput
- 10+5->order+pay
- second factor always divided by the number of replicas

# Utilization Law In Practice

How to alleviate Waiters?

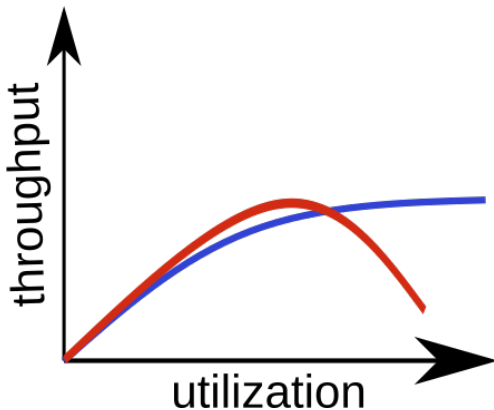
$$U = \text{throughput} \times \frac{\text{time}}{\text{replicas}}$$

- Lower the throughput
  - I.e. admit less clients
- Lower the time
  - Handle clients faster
- Increase the replicas
  - Hire more waiters



# Utilization vs Capacity

What happens when Utilization  $\approx$  Capacity?



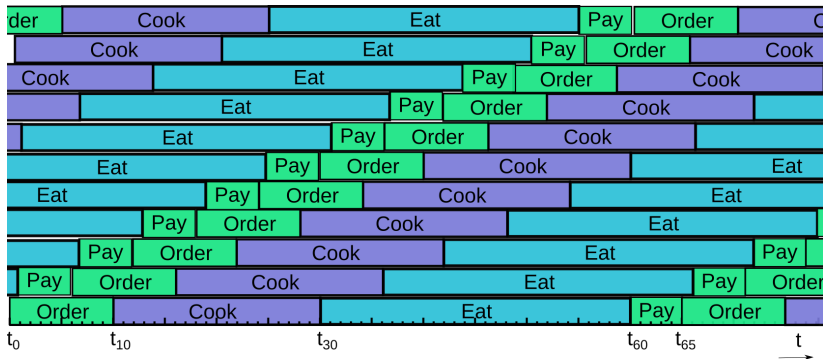
# Metrics Applied

How many clients are, on each moment, on the restaurant?

## Little's Law

The average number of pending requests in any computer system equals the average rate of arrival of requests multiplied by the average time spent by the request in the system.

# Graphical Representation



# Little's Law In Practice

- On average, how many clients are inside the restaurant?

$$c = \text{throughput} \times \text{response} = \frac{1}{6} \times 65 = 10.8$$

- What would have to be the throughput to get 15 clients inside the restaurant?

$$15 = t \times 65 \Rightarrow t = .23$$

- Homework: what about utilization?

# Metrics Applied

How many clients is each component handling?

## Forced Flow Law

The throughput through different components of a system is proportional to the number of times that component needs to handle each request.

Or

If the throughput of a component  $x$  is  $T_x$  and the throughput of the system  $S$  is  $T_s$ , each request visits  $\frac{T_x}{T_s}$  the component.

# Forced Flow Law in Practice

Component	Clients/m
Waiter	$\frac{1}{6} \times \frac{2}{3} = \frac{1}{9}$
Chef	$\frac{1}{6} \times \frac{1}{5} = \frac{1}{30}$
Table	$\frac{1}{6} \times \frac{1}{20} = \frac{1}{120}$

## ■ Notes:

- Throughput:  $\frac{1}{6}$
- $\frac{2}{3} \Rightarrow$  (Order + Pay) divided by 3 waiters
- Each waiter sees a new client on average every 9m


# Oversizing

Is the restaurant well sized?

## The $3 - \sigma$ rule

Most of the values in any distribution with a mean of  $m$  and standard deviation of  $\sigma$  lie within the range  $[m - 3\sigma, m + 3\sigma]$

# The Normal Distribution





# The 3 – $\sigma$ rule in practice

- How to avoid lines in front of my restaurant?
  - 1 Map the arrivals into a distribution
  - 2 Size the resources appropriately
- What if standard deviation is 4?
  - Expect 22 clients/h. Recalculate utilization

# Waiting Queues Theory

- So far clients were "well behaved"
  - Arrived at well-known and properly defined intervals
- What if they don't?

# Waiting Queues

- Requests don't arrive equally spaced
  - Instead they arrive with some probability  $p$  within time interval  $t$
- Popular distributions
  - Deterministic
  - General
  - Poisson

# Poisson Distribution

$$f(n, \lambda) = \frac{\lambda^n e^{-\lambda}}{n!}$$

- Probability of receiving  $n$  requests when the expected are  $\lambda$
- The probability of not receiving a request on time interval  $t$  is given by  $e^{-\lambda * t}$
- Probability of getting a request is independent of the remaining requests

# Waiting Queue Notation

- A/B/C
- Where
  - A entry distribution
  - B service distribution
  - C number of replicas
- Typical distributions
  - M Poisson
  - D Deterministic
  - G General

# Applications

With  $\lambda$  the parameter of the first distribution and  $\mu$  the parameter of the second

## M/M/1

Arrival Rate  $\lambda$

Service rate  $\mu$

Utilization  $\rho = \frac{\lambda}{\mu}$

■ If  $\rho < 1$

Average Delay  $\frac{1}{\mu(1-\rho)}$

Pending Requests  $\frac{\rho}{1-\rho}$

# Applications

With  $\lambda$  the parameter of the first distribution and  $\mu$  the parameter of the second

## M/D/1

Arrival rate  $\lambda$

Service rate  $\mu$  mas com  $\mu$  constante e fixo

Utilization  $\rho = \frac{\lambda}{\mu}$   
If  $\rho < 1$

Average Delay  $\frac{1}{2(\mu-\lambda)} + \frac{1}{2\mu}$

Pending Requests  $1 + \frac{\rho^2}{2(1-\rho)}$

# A few things on resilience

- Buzzwords
- Understanding/estimating resilience



# Buzzwords related with resilience


**MTBF** Mean Time Between Failures. The predicted time between failures of some hardware component

**N 9's** The proportion of time the system is up (when needed)

- 99.999% (5 nines): 5.26m unavailable per year
- 99.99% (4 nines): 52m unavailable per year

# Serial (alt A)

Let  $P_p$  and  $P_q$  the probability of two components (respect.)  $p$  and  $q$  be **correct**




The system will be correct if  $p$  **AND**  $q$  are both correct.  
Numerically:

$$P_p \times P_q$$

# Serial (alt B)

Let  $P_p$  and  $P_q$  the probability of two components (respect.)  $p$  and  $q$  be **incorrect**




The system will be correct if  $p$  **AND**  $q$  are both correct.

Numerically:

$$(1 - P_p) \times (1 - P_q)$$

# Parallel (alt A)


Let  $P_p$  and  $P_q$  the probability of two components (respect.)  $p$  and  $q$  be **correct**



The system will be correct if ( $p$  **AND**  $q$  are correct) **OR** ( $p$  is correct **AND**  $q$  not) **OR** ( $p$  is not correct **AND**  $q$  is correct)  $\rightarrow$   
 $(P_p \times P_q) + [P_p \times (1 - P_q)] + [(1 - P_p) \times P_q]$


# Parallel (alt B)

Let  $P_p$  and  $P_q$  the probability of two components (respect.)  $p$  and  $q$  be **correct**



The system will be **incorrect** when  $p$  **AND**  $q$  are incorrect  $\rightarrow (1 - P_p) \times (1 - P_q)$ . It will be **correct** otherwise  $\rightarrow 1 - [(1 - P_p) \times (1 - P_q)]$

# Blocks



Handle each block in separate. Compose blocks in serial.

# Wrap Up

- Learn the metrics
  - Throughput
  - Utilization
  - Response time
- Understand the dependencies between them
- Be able to justify that parallel is more resilient than serial